



Bundesministerium  
für Bildung  
und Forschung

# ATLAS Job management

## ATLAS-D Physics Meeting Heidelberg 2016

### Welcome package in Göttingen

Gen Kawamura

II.Physikalisches Institut, Universität Göttingen

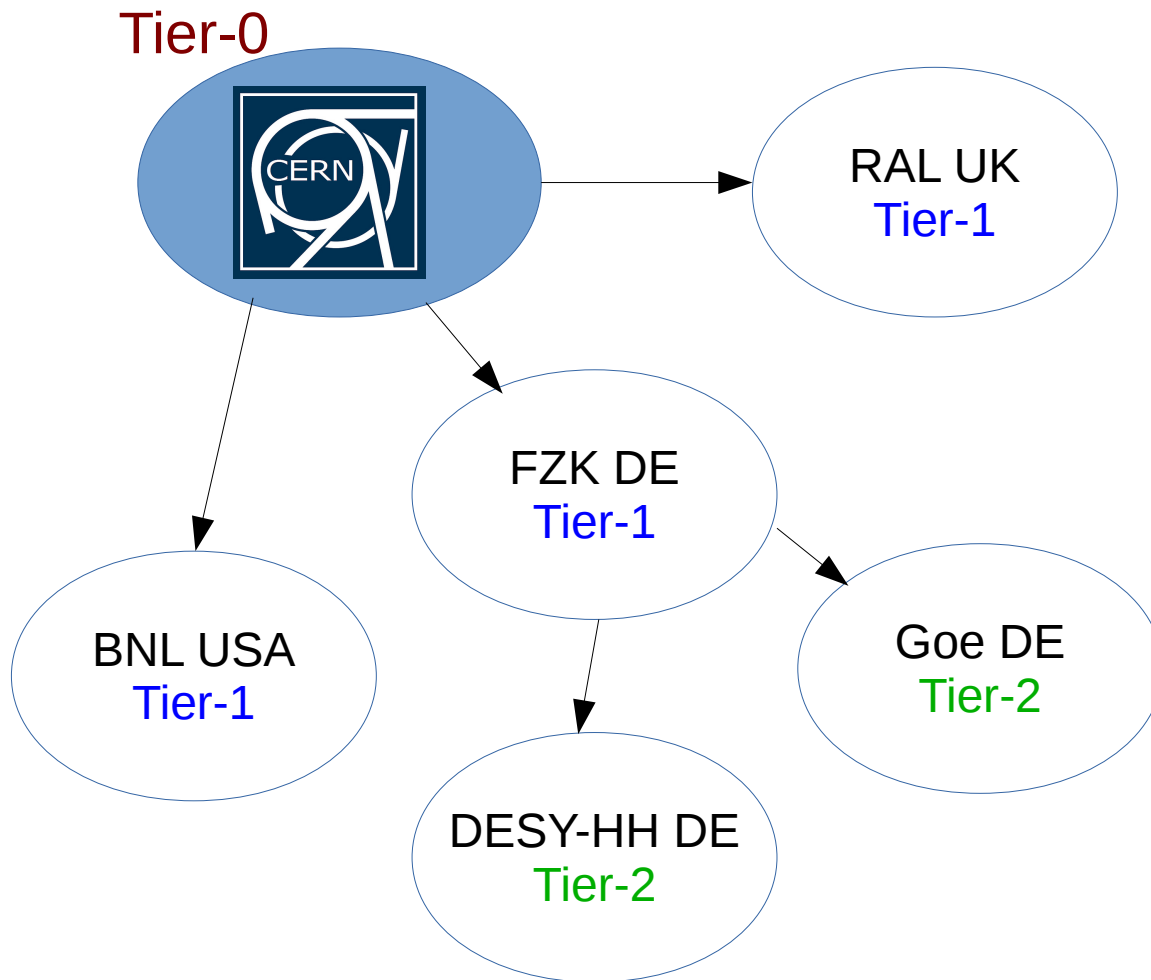
# Overview

- Your Grid environments
  - Lxplus at CERN
  - VMs at Heidelberg
  - Grid UI with Docker in your Linux Box
- Introduction to Grid computing
  - Concepts
  - Certificate Authorities and VOMS
  - Setup CVMFS
  - ATLAS Grid computing & WLCG Resources
  - Grid job and data
  - ATLAS Resources
  - Grid user interface (CLI) and CVMFS
  - BigPanda Monitoring
- Rucio (ATLAS data management system)
  - Basic concept
  - SetupRucio
  - RSE expressions
- RucioUI (WebUI)
  - RucioUI
- Links and references

# ATLAS Grid computing and WLCG resources - 1

- LHC multi-tier structure

- WLCG = Worldwide LHC Computing Grid



**Tier-0:** Raw data, Data store (in tape), Pre-processing, Reconstruction

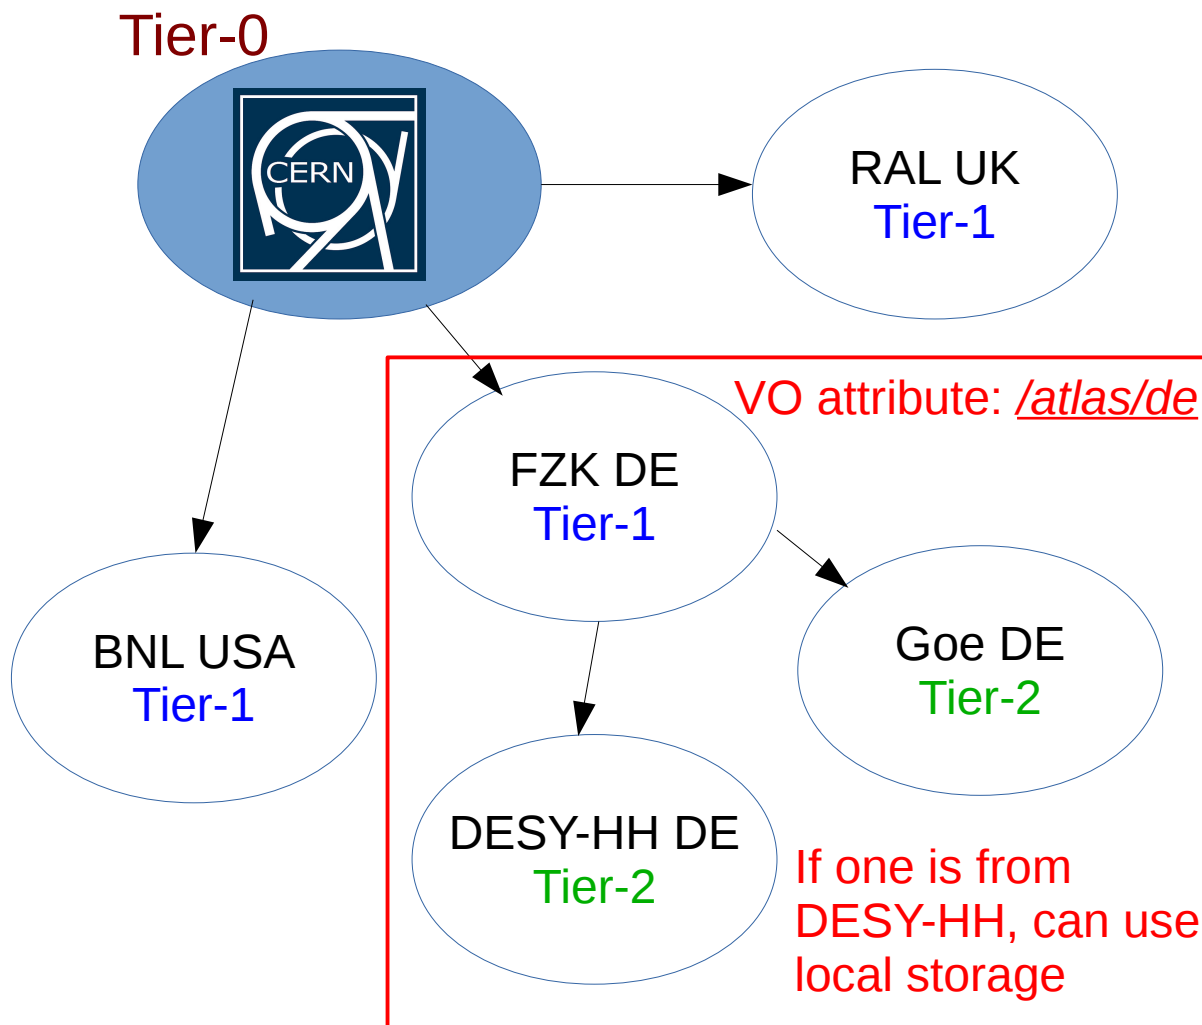
**Tier-1:** National LCG-Centre, faster network connections and larger storage spaces (e.g. Tape), MC production, user analysis, etc.

**Tier-2:** University or Facility level computing sites. MC production, user analysis, etc.

# ATLAS Grid computing and WLCG resources - 1

- LHC multi-tier structure

- WLCG = Worldwide LHC Computing Grid



**Tier-0:** Raw data, Data store (in tape), Pre-processing, Reconstruction

**Tier-1:** National LCG-Centre, faster network connections and larger storage spaces (e.g. Tape), MC production, user analysis, etc.

**Tier-2:** University or Facility level computing sites. MC production, user analysis, etc.

# Hands-on exercise

## Using ATLAS client tools

- First “Hello world” job by PanDA client

### **## PanDA client**

```
lsetup panda
```

### **## Make a Python script**

```
cat hello_world.py
```

```
#!/usr/bin/python  
print "Hello world!"
```

```
chmod 755 hello_world.py  
./hello_world.py
```

```
Hello world!
```

### **## Submitting a prun job**

```
prun --outDS user.$USER.pruntest.$$ --exec hello_world.py
```

```
INFO : gathering files under /home/gen/tmp/for_new_comer  
INFO : upload source files  
INFO : submit  
INFO : succeeded. new jediTaskID=5107461
```

### **## Submitting 5 prun jobs**

```
prun --outDS user.$USER.pruntest.$$ --exec hello_world.py -nJobs=5
```

# Hands-on exercise using ATLAS client tools

- Athena job by PanDA client
  - pathena
    - This topic would be explained in Athena session

```
## For example, you can seamlessly run Athena code on Grid
pathena DPD_topOptions.py
-inDS=data11_7TeV.00189090.physics_Egamma.merge.AOD.f403_m975
--outDS=user.wash.tutorial.dpd
```

# BigPanda Monitoring

- Next generation monitoring tool developed within the BigPanDA project
- Drop down menu with access to numerous views of job monitoring data
- Tool heavily used by ADC experts and shifters to track production activities
- Here we are more interested in monitoring our own jobs



# BigPanda Search Options

Search	
PanDA job ID or name	<input type="text"/> <input type="button" value="Submit"/>
Batch ID	<input type="text"/> <input type="button" value="Submit"/>
Task ID	<input type="text"/> <input type="button" value="Submit"/>
Task name	<input type="text"/> <input type="button" value="Submit"/>
Request ID	<input type="text"/> <input type="button" value="Submit"/>
Tasks for Request IDs	from <input type="text"/> to <input type="text"/> <input type="button" value="Submit"/>
Jobs for Request IDs	from <input type="text"/> to <input type="text"/> <input type="button" value="Submit"/>
Tasks for campaign	<input type="text"/> <input type="button" value="Submit"/>
Tasks for project	<input type="text"/> <input type="button" value="Submit"/>
Tasks for tag	<input type="text"/> <input type="button" value="Submit"/>
Tasks for stream	<input type="text"/> <input type="button" value="Submit"/>
Site	<input type="text"/> <input type="button" value="Submit"/>
User	<input type="text"/> <input type="button" value="Submit"/>
Dataset	<input type="text"/> <input type="button" value="Submit"/>
File	<input type="text"/> <input type="button" value="Submit"/>

- Search boxes available on BigPanda front page to get specific information on tasks, jobs and user activities
- The range of displayed information can be changed by modifying or adding the URL parameters

Enter our example  
job Task ID here

The load on the Panda  
DB will increase with  
high N

```
&limit=N  
&display_limit=N  
&hours=N  
&days=N  
&earlierthan=Nhours  
&earlierthandays=Ndays  
&date_from=2014-10-01  
&date_to=2014-10-03
```

## Example URL Modifiers

- Advanced usage: monitoring data can be retrieved in *json* format to be able results to be processed programmatically

➤ <http://bigpanda.cern.ch/help/>



# User View

- The user view displays your recent activities
- Useful starting point to check your jobs
- Links on this page lead to more information on an individual tasks and associated jobs

ATLAS PanDA Dash Tasks Jobs Errors Users Sites Incidents Search

PanDA user andrew washbrook, last 72 hours, limit 5000 per table. Params: limit=5000 user=washbrook

User page for andrew washbrook Please read (and keep an eye on) the help section below to familiarize yourself with using the monitor in JEDI based analysis.

Stats from the user DB (updated every few hours)

User	Jobs in last week	Latest job in the DB stats	Personal CPU-hrs - 1 day	Personal CPU-hrs - 7 days	Group CPU-hrs - 1 day	Group CPU-hrs - 7 days
andrew washbrook	87	2013-10-12 19:33	0.0	0.0	0.0	0.0

Jump to recent tasks jobsets jobs

Task attribute summary, 1 tasks

corecount (1)	1 (1)
processingtype (1)	panda-client-0.5.42-jedi-run (1)
reqid (1)	3 (1)
status (1)	done (1)
taskpriority (1)	1000 (1)
tasktype (1)	analy (1)
transpath (1)	runGen-00-00-02 (1)
username (1)	andrew washbrook (1)

Grid resources used

Task Summary

Replace with your username

➤ <http://bigpanda.cern.ch/user/?user=washbrook>

1 tasks, sorted by jedid taskid Show in task list page

ID Parent	Jobset	Task name TaskType/ProcessingType Campaign Group User Logged status	Task status Nfiles	Input files finish% fail% Nfinish Nfail	Modified	State changed	Priority
4617520	3	user.wash.prntest/ anal/panda-client-0.5.42-jedi-run andrew washbrook	done 1	100% 1	2015-01-09 10:15	01-09 10:15	1000

Task Details

Recent jobsets Note that with JEDI you should pay more attention to tasks than jobsets! See the help section below.

Jobset ID	nJobs	Modification times	Priorities	defined	waiting	assigned	activated	sent	starting	running	holding	transferring	finished	failed	cancelled
1	2	2015-01-06 14:28 to 01-06 14:36	1000 to 2000	0	0	0	0	0	0	0	0	0	2	0	0

# Example Job Monitoring

➤ <http://bigpanda.cern.ch/task/4617520/>

Task 4617520: user.wash.pruntest/											
Task ID	Jobset	Type	WorkingGroup	User	Task status	Ninputfiles   finished   failed	Created	Modified	Cores	Priority	Parent
4617520	3	analy		andrew washbrook	done	1   1 (100%)	2015-01-06 14:22	01-09 10:15	1	1000	

<b>View:</b>	job list (access to job details and logs)	child tasks	prodsys task page	brokerage logger	JEDI action logger	error summary
--------------	---	-------------	-------------------	------------------	--------------------	---------------



➤ <http://bigpanda.cern.ch/job?pandaaid=2351340829>

Where the job was run



PandaID	Owner	Request Task ID	Status	Created	Time to start d:h:m:s	Duration d:h:m:s	Modified	Cloud Site	Priority
2351340829	andrew washbrook	1 4617520	finished	2015-01-06 14:23	0:0:05:05	0:0:01:13	01-06 14:36	IT ANALY_INFN-FRASCATI	1000

**Job name:** user.wash.pruntest/ **type:** panda-client-0.5.42-jedi-run **transformation:** runGen-00-00-02

**Datasets:** Out: user.wash.pruntest.log.14345588

<b>Job information</b>	logfiles	pilot job stdout, stderr, batch log	logger records	child jobs
------------------------	----------	-------------------------------------	----------------	------------

<b>Jobset 1</b>
-----------------



# Job Log Files

- [http://bigpanda.cern.ch/filebrowser/?guid=b2c3305e-5581-4062-9676-92205d72018a&lfname=user.wash.prntest.log.4617520.000001.log.tgz&site=ANALY\\_INFN-FRASCATI&scope=user.wash](http://bigpanda.cern.ch/filebrowser/?guid=b2c3305e-5581-4062-9676-92205d72018a&lfname=user.wash.prntest.log.4617520.000001.log.tgz&site=ANALY_INFN-FRASCATI&scope=user.wash)
- Annotations:
- Globally Unique Identifier (points to `guid=b2c3305e-5581-4062-9676-92205d72018a`)
  - Logical File Name (points to `lfname=user.wash.prntest.log.4617520.000001.log.tgz`)

## File listing

Modification time	Size [B]	File
Tue Jan 6 14:28:57 2015	0	<a href="#">athena_stderr.txt</a>
Tue Jan 6 14:28:57 2015	23820	<a href="#">athena_stdout.txt</a>
Tue Jan 6 14:28:57 2015	730	<a href="#">job_setup.sh</a>
Tue Jan 6 14:28:52 2015	188	<a href="#">matched_replicas.json</a>
Tue Jan 6 14:28:57 2015	527	<a href="#">metadata-2351340829.xml</a>
Tue Jan 6 14:28:49 2015	1391	<a href="#">pandaJobData.out</a>
Tue Jan 6 14:28:57 2015	0	<a href="#">pandatracerlog.txt</a>
Tue Jan 6 14:28:49 2015	0	<a href="#">pilot.stderr</a>
Tue Jan 6 14:28:50 2015	77	<a href="#">PILOT_INITDIR</a>
Tue Jan 6 14:28:57 2015	97	<a href="#">Pilot_VmPeak.txt</a>
Tue Jan 6 14:28:51 2015	280	<a href="#">pilotlog.out</a>
Tue Jan 6 14:29:52 2015	83922	<a href="#">pilotlog.txt</a>

- From an individual job view you can get access to log files for further diagnosis

Output and error of job payload

Pilot logs show the staging of the job at the site

# Task Monitoring

- The new task is visible from my user view page in a *registered* state

2 tasks, sorted by jedidaskid <a href="#">Show in task list page</a>									
ID Parent	Jobset	Task name TaskType/ProcessingType Campaign Group User Logged status	Task status Nfiles	Input files finish% fail% Nfinish Nfail	Modified	State changed	Priority		
4635475	6	user.wash.tutorial.dpd/ anal/panda-client-0.5.42-jedi-athena andrew washbrook	registered 29		2015-01-09 15:47	01-09 15:47	1000		
4617520	5	user.wash.pruntest/ anal/panda-client-0.5.42-jedi-run andrew washbrook	done 1	100% 1	2015-01-09 14:44	01-09 14:44	1000		

Task 4635475: user.wash.tutorial.dpd/									
Task ID	Jobset	Type	WorkingGroup	User	Task status	Detailed JEDI status	Ninputfiles   finished   failed	Created	Modified
4635475	6	analy		andrew washbrook	submitting	scouting	29   0 (0%)	2015-01-09 15:47	01-09 15:48

States of jobs in this task (merge jobs excluded) <a href="#">Show all jobs</a>														
defined	waiting	pending	assigned	throttled	activated	sent	starting	running	holding	transferring	finished	failed	cancelled	merging
11														

Jump to [job parameters](#), [task parameters](#)

View:	<a href="#">job list (access to job details and logs)</a>	<a href="#">child tasks</a>	<a href="#">prodsys task page</a>	<a href="#">brokerage logger</a>	<a href="#">JEDI action logger</a>	<a href="#">error summary</a>
-------	---	-----------------------------	-----------------------------------	----------------------------------	------------------------------------	-------------------------------

Input containers
<a href="#">data11_7TeV.00189090.physics_Egamma.merge.AOD.f403_m975</a>

Output containers
<a href="#">user.wash.tutorial.dpd_MyFirstD3PD/</a>
<a href="#">user.wash.tutorial.dpd.log/</a>

➤ <http://bigpanda.cern.ch/user/?user=washbrook>

# Task Monitoring (2)

Job attribute summary		Sort by <b>count</b> , alpha
atlasrelease (1)	Atlas-17.2.7 (11)	
attemptnr (1)	1 (10)	
cloud (1)	DE (11)	
computingsite (1)	ANALY_FMPHi-UNIBA (11)	
destinationse (1)	ANALY_FMPHi-UNIBA (1)	
homepackage (1)	AnalysisTransforms-AtlasPhysics_17.2.7.4.1 (11)	
inputfileproject (1)	data11_7TeV (10)	
inputfiletype (1)	AOD (10)	
jeditaskid (1)	4635475 (11)	
jobsetid (1)	6 (11)	
jobstatus (2)	defined (10) running (1)	
outputfiletype (1)	? (11)	
priorityrange (3)	900:999 (6) 1000:1099 (4) 2000:2099 (1)	
processingtype (1)	panda-client-0.5.42-jedi-athena (11)	
prodsourcelabel (2)	panda (1) user (10)	
produsername (1)	andrew washbrook (11)	
reqid (1)	6 (11)	
specialhandling (1)	ddm:rucio (11)	
transformation (2)	buildJob-00-00-03 (1) runAthena-00-00-12 (10)	

➤ <http://bigpanda.cern.ch/task/4635475/>

- Switching to the Task view we can now see one job is running at the computing site ANALY\_FMPHi-UNIBA
- Can follow individual progression of running jobs by clicking on Job Status

Job list    Sort by PandalID, ascending mod time, descending mod time, priority, attemptnr, duration										
PanDA ID Attempt#	Owner Group	Request Task ID	Transformation	Status	Created	Time to start d:h:m:s	Duration d:h:m:s	Mod	Cloud Site	Priority
2354305514 Attempt 0	andrew washbrook	6 4635475	buildJob-00-00-03	running	2015-01-09 15:48	0:0:01:16	0:0:01:19	01-09 15:49	DE ANALY_FMPHl-UNIBA	2000
	Job name: user.wash.tutorial.dpd/ #0									
	Datasets:    Out: panda.0109154800.984205.lib_4635475									

# Task Monitoring (3)

States of jobs in this task (merge jobs excluded) <a href="#">Show all jobs</a>														
defined	waiting	pending	assigned	throttled	activated	sent	starting	running	holding	transferring	finished	failed	cancelled	merging
								10			1			

Job attribute summary		Sort by <b>count</b> , alpha
atlasrelease (1)	Atlas-17.2.7 (11)	
attemptnr (1)	1 (10)	
cloud (1)	DE (11)	
computingsite (1)	ANALY_FMPHi-UNIBA (11)	
destinationse (1)	ANALY_FMPHi-UNIBA (1)	
homepackage (1)	AnalysisTransforms-AtlasPhysics_17.2.7.4.1 (11)	
inputfileproject (1)	data11_7TeV (10)	
inputfiletype (1)	AOD (10)	
jeditaskid (1)	4635475 (11)	
jobsetid (1)	6 (11)	
jobstatus (2)	finished (1) running (10)	
outputfiletype (1)	? (11)	
priorityrange (3)	900:999 (4) 1000:1099 (6) 2000:2099 (1)	
processingtype (1)	panda-client-0.5.42-jedi-athena (11)	
prodsourcelabel (2)	panda (1) user (10)	
produsername (1)	andrew washbrook (11)	
reqid (1)	6 (11)	
specialhandling (1)	ddm:rucio (11)	
transformation (2)	buildJob-00-00-03 (1) runAthena-00-00-12 (10)	

➤ <http://bigpanda.cern.ch/task/4635475/>

- Checking some time later there are 10 jobs in a *running* state with 1 job *finished*

Job list										
Sort by PandaID, ascending mod time, descending mod time, priority, attemptnr, duration										
PanDA ID Attempt#	Owner Group	Request Task ID	Transformation	Status	Created	Time to start d:h:m:s	Duration d:h:m:s	Mod	Cloud Site	Priority
2354305524 Attempt 1	andrew washbrook	6 4635475	runAthena-00-00-12	running	2015-01-09 15:48	0:0:15:42	0:0:05:08	01-09 16:03	DE ANALY_FMPHI-UNIBA	999
	Job name: user.wash.tutorial.dpd/ #1									
	Datasets: In: data11_7TeV:data11_7TeV.00189090.physics_Egamma.merge.AOD.f403_m975 Out: user.wash.tutorial.dpd_MyFirstD3PD.14603408									
2354305523 Attempt 1	andrew washbrook	6 4635475	runAthena-00-00-12	running	2015-01-09 15:48	0:0:15:42	0:0:05:08	01-09 16:03	DE ANALY_FMPHI-UNIBA	1000
	Job name: user.wash.tutorial.dpd/ #1									
	Datasets: In: data11_7TeV:data11_7TeV.00189090.physics_Egamma.merge.AOD.f403_m975 Out: user.wash.tutorial.dpd_MyFirstD3PD.14603408									
2354305522 Attempt 1	andrew washbrook	6 4635475	runAthena-00-00-12	running	2015-01-09 15:48	0:0:15:42	0:0:05:08	01-09 16:03	DE ANALY_FMPHI-UNIBA	1000
	Job name: user.wash.tutorial.dpd/ #1									
	Datasets: In: data11_7TeV:data11_7TeV.00189090.physics_Egamma.merge.AOD.f403_m975 Out: user.wash.tutorial.dpd_MyFirstD3PD.14603408									

# Task Monitoring (4)

Task 4635475: user.wash.tutorial.dpd/

Task ID	Jobset	Type	WorkingGroup	User	Task status	Ninputfiles   finished   failed	Created	Modified	Cores	Priority	Parent
4635475	6	analy		andrew washbrook	done	29   29 (100%)	2015-01-09 15:47	01-09 17:49	1	1000	

States of jobs in this task (merge jobs excluded)

Show all jobs

defined	waiting	pending	assigned	throttled	activated	sent	starting	running	holding	transferring	finished	failed	cancelled	merging
											15			

Job attribute summary <a href="#">Sort by count, alpha</a>	
atlasrelease (1)	Atlas-17.2.7 (15)
attemptnr (2)	1 (10) 2 (3)
cloud (2)	DE (11) UK (4)
computingsite (2)	ANALY_FMPH-UNIBA (11) ANALY_SHEF_SL6 (4)
destinationse (2)	ANALY_FMPH-UNIBA (1) ANALY_SHEF_SL6 (1)
homepackage (1)	AnalysisTransforms-AtlasPhysics_17.2.7.4.1 (15)
inputfileproject (1)	data11_7TeV (13)
inputfiletype (1)	AOD (13)
jeditaskid (1)	4635475 (15)
jobsetid (1)	6 (15)
jobstatus (1)	finished (15)
outputfiletype (1)	? (15)
priorityrange (3)	900:999 (4) 1000:1099 (9) 2000:2099 (2)
processingtype (1)	panda-client-0.5.42-jedi-athena (15)
prodsourcelabel (2)	panda (2) user (13)
produsername (1)	andrew washbrook (15)
reqid (1)	6 (15)
specialhandling (1)	ddm:rucio (15)
transformation (2)	buildJob-00-00-03 (2) runAthena-00-00-12 (13)

➤ <http://bigpanda.cern.ch/task/4635475/>

- And then later all the jobs have finished and the task status is *done*
- Note that 4 of the 15 jobs were sent to a computing site in a different cloud (UK)
- Do not expect all your jobs to run at the same computing site

# Links and references

- RucioUI
  - <https://rucio-ui.cern.ch/>
- Rucio Documentation
  - <http://rucio.cern.ch/index.html>
- Software twiki tutorial
  - <https://twiki.cern.ch/twiki/bin/viewauth/AtlasComputing/SoftwareTutorialGettingDatasets>
- Athena Docker setup
  - <https://twiki.cern.ch/twiki/bin/viewauth/AtlasComputing/AthenaMacDockerSetup>
- Docker container for CVMFS
  - <https://github.com/sbinet/docker-containers/tree/master/cvmfs-atlas>
- Binet, Sébastien, and Ben Couturier. "*docker & HEP: Containerization of applications for development, distribution and preservation.*" Journal of Physics: Conference Series. Vol. 664. No. 2. IOP Publishing, 2015.
  - <http://iopscience.iop.org/article/10.1088/1742-6596/664/2/022007/meta>
- ATLAS-D meeting 2015 Rucio Tutorial, Thomas Beermann
- Monitoring Your Grid Jobs, Andrew Washbrook University of Edinburgh, ATLAS Software & Computing Tutorials 14th January 2015 PUC, Chile
- Athena Mac Docker
  - <https://twiki.cern.ch/twiki/bin/viewauth/AtlasComputing/AthenaMacDockerSetup>
- Software tutorial using Grid
  - <https://twiki.cern.ch/twiki/bin/viewauth/AtlasComputing/SoftwareTutorialUsingTheGrid>